



The dark art of performance tuning

or how to become a perf hero



Agenda

- Introduction to benchmarking and profiling
- Profiling metrics and tools - O/S level
- JVM primer
- Profiling metrics and tools - JVM level
- JVM Benchmarking
 - pitfalls
 - JMH





Benchmarking vs profiling

- Benchmarking
 - measuring different solutions to the same problem, e.g. sorting
- Profiling
 - runtime program analysis that yields execution information, e.g.
 - frequency and duration of particular function calls
 - memory allocation rate
- Usually go hand in hand
 - benchmark -> profile -> make a change -> benchmark

Performance analysis - why

- expectations vs reality
- new client requirements
 - SLAs to meet, e.g. 10k ops/s
 - new features
- marketing
- gathering knowledge to make conscious decisions
- The First Rule of Program Optimization: *Don't do it.*
The Second Rule of Program Optimization (for experts only!): *Don't do it yet*
 - prefer readability to complicated but more performant code in MOST of the cases
 - optimize hotspots and bottlenecks
- fun & profit

Latency numbers every programmer should know



Benchmarks amplify all the effects visible at the same scale

- **kilo**: > 1000 s, Linpack
- **_____**: 1...1000 s, SPECjvm2008, SPECjbb2013
- **milli**: 1...1000 ms, SPECjvm98, SPECjbb2005 not really hard
- **micro**: 1...1000 us, single webapp request challenging
- **nano**: 1...1000 ns, single operations damned beasts
- **pico**: 1...1000 ps, pipelining ...

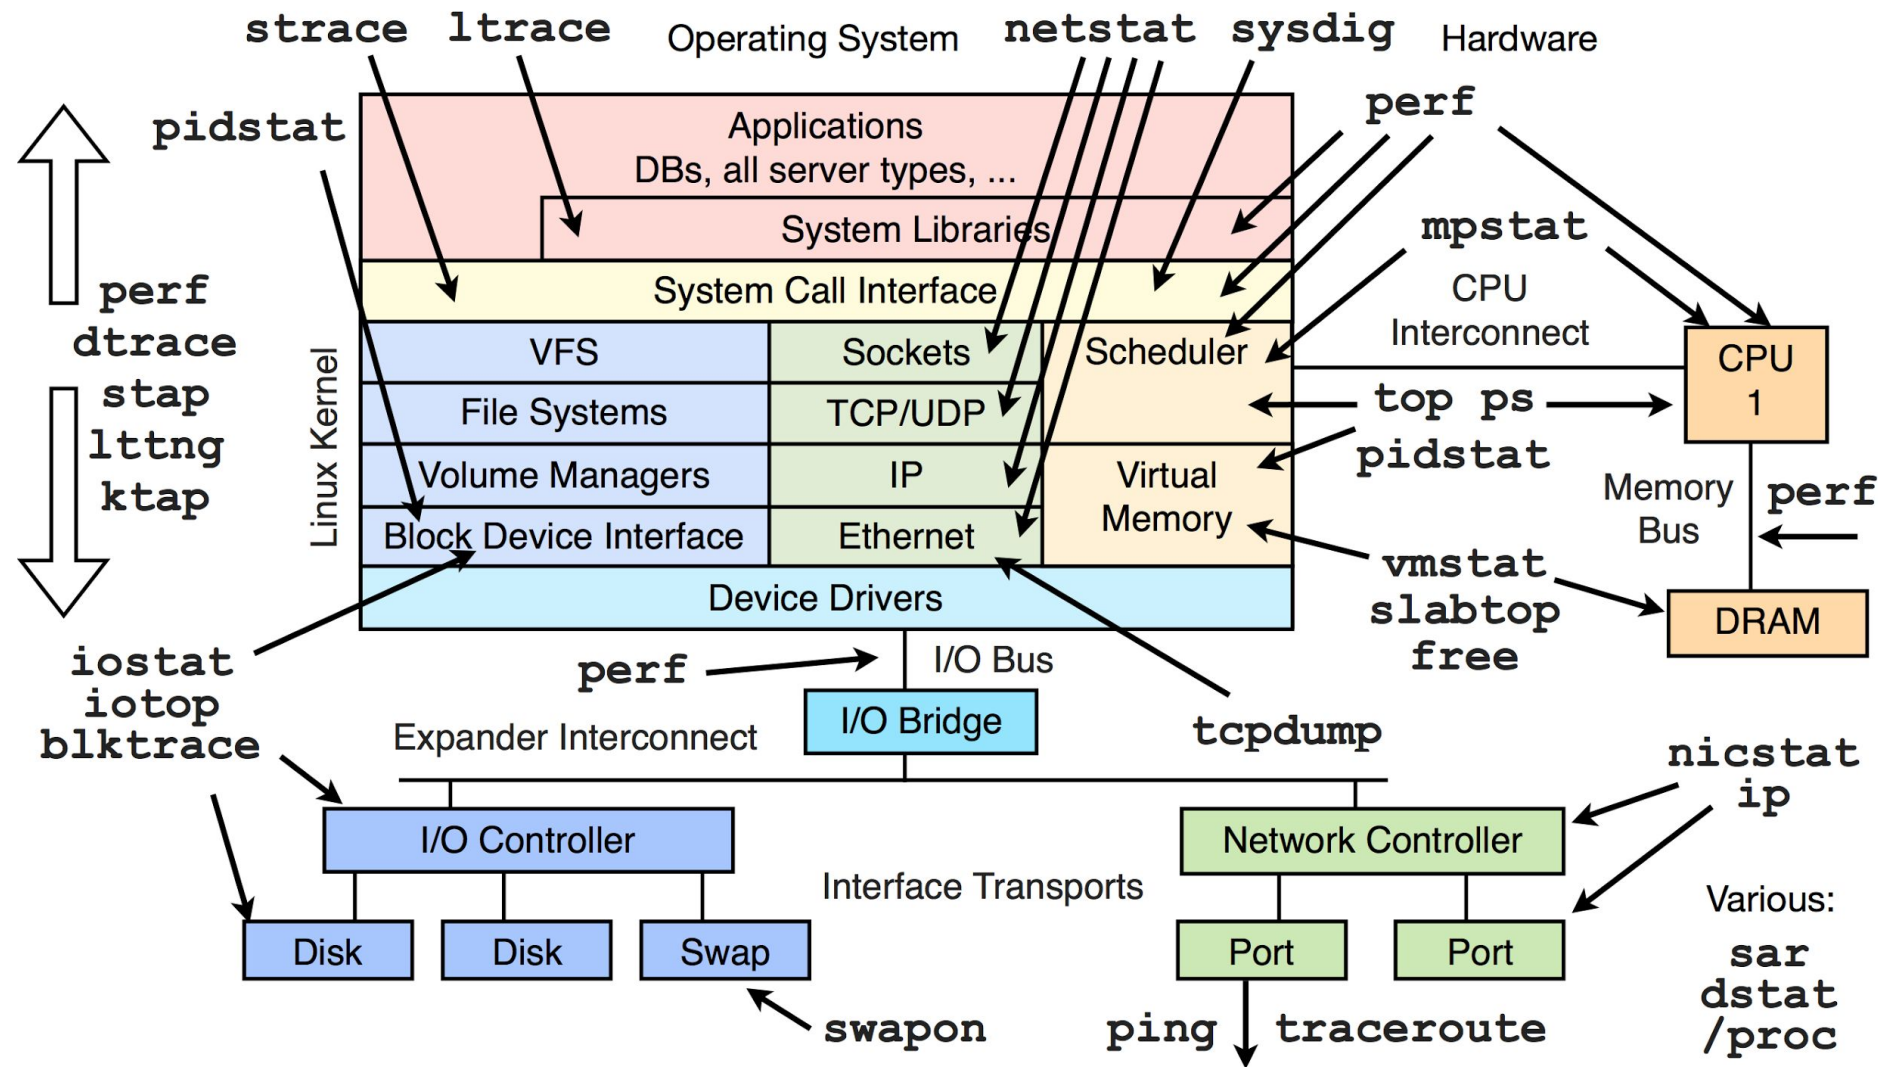
Optimization quiz

Empty benchmark, system reports 4 online CPUs

Threads	Ops/nsec	Scale
1	3.06 ± 0.10	
2	5.72 ± 0.10	1.87 ± 0.03
4	5.87 ± 0.02	1.91 ± 0.03

Software is *abstract*
hardware is **real**

O/S level



O/S level (Linux)

- top/htop/ps
 - overall state of the system:
 - list of processes and threads
 - cpu usage
 - memory/swap usage
 - ...
- vmstat
 - overall state of the system:
 - memory/swap usage
 - basic i/o information
 - interrupts/context switches
 - cpu usage - system, user, idle, waiting for I/O,

```
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r  b   swpd   free   buff  cache   si   so    bi    bo    in   cs  us  sy  id  wa  st
3  0       3   1257   137   1390    0    0     0     0   906 2064  6  1 94  0  0
1  0       3   1257   137   1390    0    0     0     0  1074 2537  3  1 96  0  0
```

- `vmstat -S m`
 - r - processes waiting for run time
 - b - processes blocked, waiting on resources
 - swpd - swapped memory
 - free - free memory
 - buff/cache - memory used as buffers/cache
 - si/so - swap in/swap out in pages
 - bi/bo - blocks received/sent from/to block device
 - in - interrupts
 - cs - context switches
 - us/sy/id/wa/st - CPU time

O/S level (Linux)

- iostat/iotop
 - I/O requests per second
 - I/O requests completed
 - average size of the requests
 - average wait time of requests to be completed (queue + service)
 - ...
- pidstat
 - various statistics for linux tasks
 - I/O, page faults and memory utilization, context switches
- nload/iftop
 - networking statistics

Total DISK READ :		0.00 B/s	Total DISK WRITE :		2.46 M/s		
Actual DISK READ:		0.00 B/s	Actual DISK WRITE:		0.00 B/s		
TID	PRI	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
4168	be/4	cassandr	0.00 B/s	143.61 K/s	0.00 %	0.24 %	java -ea -javaagent:/usr/share/cassandra/1
3973	be/4	cassandr	0.00 B/s	191.48 K/s	0.00 %	0.17 %	java -ea -javaagent:/usr/share/cassandra/1
5859	be/4	cassandr	0.00 B/s	62.60 K/s	0.00 %	0.09 %	java -ea -javaagent:/usr/share/cassandra/1
3972	be/4	cassandr	0.00 B/s	228.31 K/s	0.00 %	0.04 %	java -ea -javaagent:/usr/share/cassandra/1
3970	be/4	cassandr	0.00 B/s	301.96 K/s	0.00 %	0.00 %	java -ea -javaagent:/usr/share/cassandra/1
4165	be/4	cassandr	0.00 B/s	228.31 K/s	0.00 %	0.00 %	java -ea -javaagent:/usr/share/cassandra/1
4166	be/4	cassandr	0.00 B/s	180.44 K/s	0.00 %	0.00 %	java -ea -javaagent:/usr/share/cassandra/1
4167	be/4	cassandr	0.00 B/s	235.67 K/s	0.00 %	0.00 %	java -ea -javaagent:/usr/share/cassandra/1
5778	be/4	adebski	0.00 B/s	29.46 K/s	0.00 %	0.00 %	java -XX:+PrintGCDetails -Xloggc:gc-145997
5779	be/4	adebski	0.00 B/s	33.14 K/s	0.00 %	0.00 %	java -XX:+PrintGCDetails -Xloggc:gc-145997
5780	be/4	adebski	0.00 B/s	29.46 K/s	0.00 %	0.00 %	java -XX:+PrintGCDetails -Xloggc:gc-145997
5791	be/4	adebski	0.00 B/s	29.46 K/s	0.00 %	0.00 %	java -XX:+PrintGCDetails -Xloggc:gc-145997
5793	be/4	adebski	0.00 B/s	7.36 K/s	0.00 %	0.00 %	java -XX:+PrintGCDetails -Xloggc:gc-145997

- iotop --only
 - SWAPIN - percentage of time the thread spent while swapping in
 - IO - percentage of time the thread spent while waiting on I/O

Perf events - Linux

- lightweight profiling solution for Linux systems
- included in the kernel
- based on notion of events
 - cache misses
 - branch mispredictions
 - page faults
 - context switches
 - cpu time
 - system calls
 - ...
- disadvantage: hard to find documentation for specific events

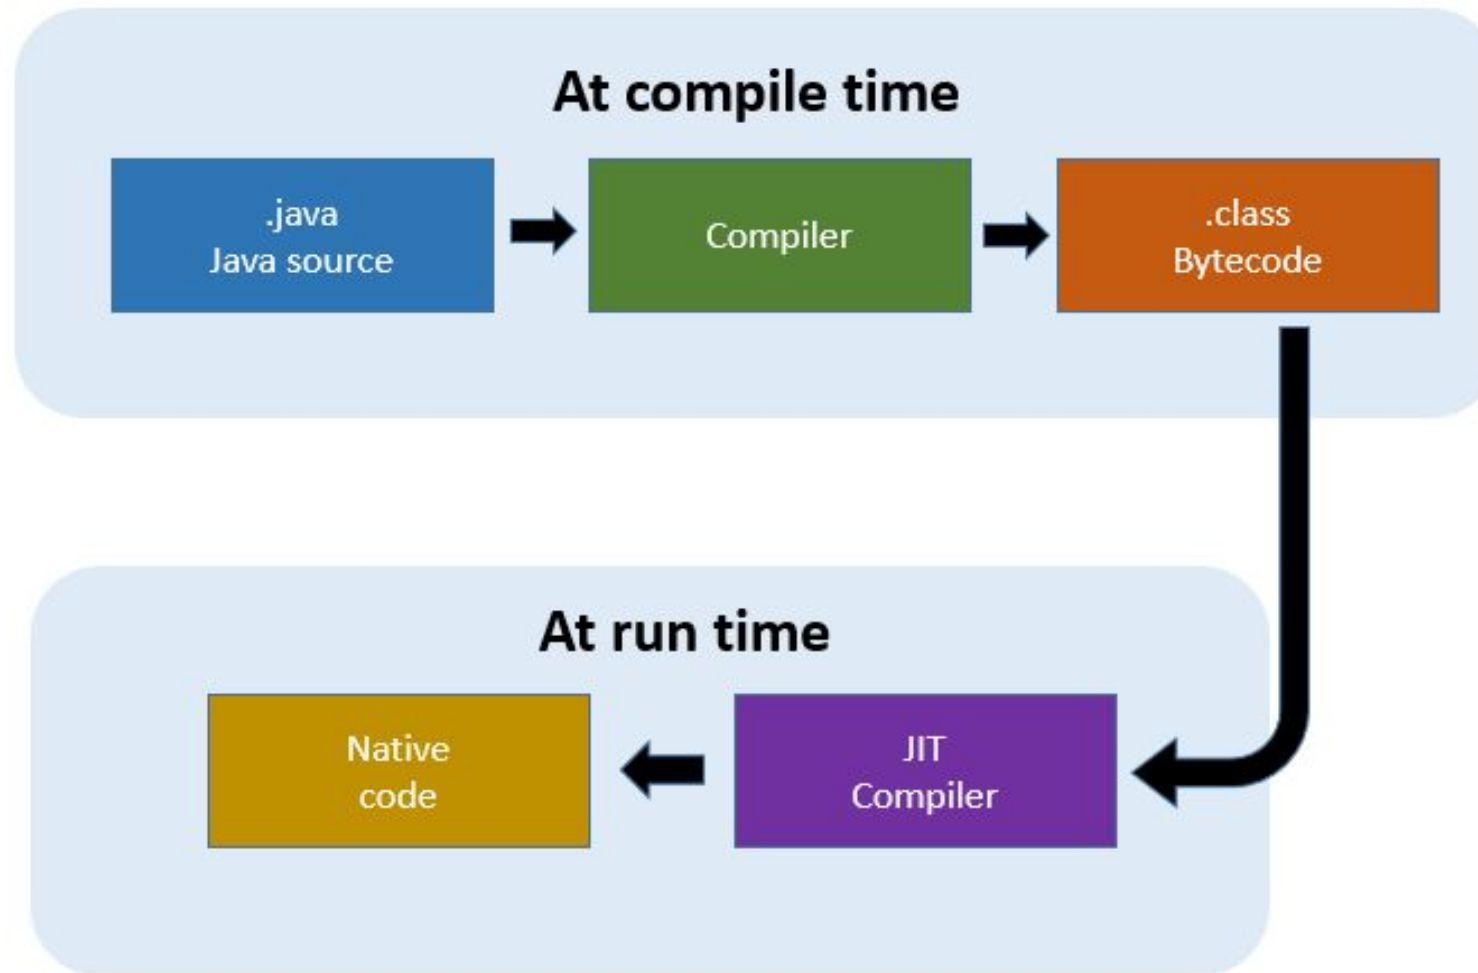


Flamegraphs

- <http://www.brendangregg.com/flamegraphs.html>
- a way to visualize profiling data, not tied to perf tool
- perf record + perf script + perl script = svg image
- problem: for stacktraces to be readable requires debug symbols
- what about languages that run in VMs:
 - java - <https://bugs.openjdk.java.net/browse/JDK-8068945>
 - ruby
 - javascript
 - ...

JVM internals - primer

JIT compiler




JIT compiler

- introduced to Sun JVM around java 1.2/1.3
- almost no optimizations during initial compilation (javac), generated bytecode mostly reflects .java code
- at the beginning bytecode is interpreted, only hot methods are compiled to native code
- advantages:
 - runtime information can lead to interesting optimizations, e.g. inlining virtual method call because only single class with given interface is currently loaded, removing null checks
 - leveraging platform specific optimizations during runtime
 - code written and compiled for java X can benefit from future performance improvements in java X + 1

JIT compiler

- disadvantages
 - initial performance is very low (JVM warmup)
 - compilation and optimization costs
 - unpredictable
 - optimistic assumptions can fail - code returns to being interpreted and awaits second JIT compilation

Memory layout




<http://www.oracle.com/webfolder/technetwork/tutorials/obe/java/gc01/index.html>

Garbage collection

- cheap allocation, paying the price later during GC
- STW (stop the world) pause
- more than single algorithm available
 - single threaded STW
 - multi threaded STW
 - CMS - only old generation, “mostly concurrent”
 - G1 - both young and old generation, “mostly concurrent”
- many tuning parameters:
 - usually best to use the defaults and rely on JVM runtime analysis
 - using arcane options can lead to very fragile performance
 - single JVM - single GC - multiple allocation profiles and requirements

Warmup

JDK8b83 + nashorn 2013-04-08, Octane:DeltaBlueBench performance over iterations



<http://shipilev.net/talks/devoxx-Nov2013-benchmarking.pdf>

Profiling metrics and tools - JVM level

Command line tools

- bin directory has 49 executables, some of them are:
 - **jstat** - GC and JIT statistics
 - **jstack** - stack trace of all threads, deadlock detection
 - **jps** - list of java processes, prints JVM options
 - **jinfo** - more detailed information about java process
 - **jmap** - heap histogram, heap dump
- available on each system with JDK

Command line options

- GC logs
 - `-XX:+PrintGC -XX:+PrintGCDetails -XX:+PrintGCTimeStamps -XX:+PrintGCDateStamps -Xloggc:<file>`
 - very useful when investigating GC issues
- Peeking under the JIT hood
 - add `-XX:+UnlockExperimentalVMOptions` just to be on the safe side
 - `-XX:+PrintCompilation -XX:+CITime` - basic information about compilation
 - `-XX:+PrintInlining` - outputs information what methods were inlined and what weren't and why
 - `-XX:+PrintAssembly` - shows what assembly code was generated by the JIT compiler

Profilers - VisualVM

- bundled with JDK, free of charge
- CPU information
- detailed memory information with VisualGC
- sampling (CPU/memory)
- profiling (CPU/memory)
- JMX browser
- thread information
- with a little effort can be used with remote JVMs

Profilers - Oracle Mission Control

- Everything available in VisualVM and more
- Free for development, have to pay when used in production
- Very low overhead (according to the Oracle) due to using internal and undocumented APIs
- Available since JDK 7u40

Eclipse memory analyzer

- analyzes full heap dump that can be obtained using jmap tool
- it can be used to obtain information like:
 - finding biggest objects on the heap
 - learning why specific objects are not removed by GC (GC roots)
 - listing dominating objects on the heap
- allows to inspect field values of specific objects
- may require significant memory during first pass through the heap dump but subsequent openings will be a lot faster

Microbenchmarking

What?

Micro + benchmark

performance measurement of a **very small** piece of code, something that might take μ s or ns to run

How?

Java Microbenchmarking Harness

- benchmark code generation, driven by annotations.
- generated classes and all their dependencies get packaged in an all-in-one runnable jar
- benchmark runner supporting single threaded, multi-threaded and thread groups
- pluggable profilers
- multi-language support
- reporting formats - JSON, CSV, etc.

JMH examples

```
package jmh

import org.openjdk.jmh.annotations._

/**
 * @author nuk
 */
@Warmup(iterations = 5)
@Measurement(iterations = 20)
@Fork(1)
@Threads(8)
class HelloWorld {

  @Benchmark
  def yo(): Unit = {
    // this method was intentionally left blank.
  }

}
```

JMH examples

```
@Benchmark
@OperationsPerInvocation(OpBatch)
def setAddAndRemove(bh: Blackhole, state: BenchmarkState): Unit = {
  import state.manager.executor
  val strings = (1 to OpBatch).map(_ => nextString())
  val f1 = strings.map(state.set.add(_))
  val f2 = strings.map(state.set.remove(_))
  bh.consume(Await.result(Future.sequence(f1), timeout))
  bh.consume(Await.result(Future.sequence(f2), timeout))
}
```

JMH examples

```
@State(Scope.Benchmark)
class BenchmarkState {

  val manager = new WorManager(
    new Config("127.0.0.1", 6379, 1, "jmh"),
    new ActorSystemProvider {lazy val system = ActorSystem("WorOpBenchmark")}
  )
  var wor: Wor = _
  var value: WorValue = _
  var set: WorSet = _
  var log: WorHyperLogLog = _
  var list: WorList = _


  @Setup(Level.Iteration)
  def setup(): Unit = {
    Await.result(manager.executor.flushdb(), timeout)
    import manager.executor
    wor = Await.result(manager.getOrCreateWor("values"), timeout)
    val f = Future.sequence(
      Seq(
        wor.createValue("value", "1").map(value = _),
        wor.createSet("set", "a").map(set = _),
        wor.createHyperLogLog("log", "a").map(log = _),
        wor.createList("list", "a").map(list = _)
      )
    )
    Await.result(f, timeout)
  }

  @TearDown(Level.Iteration)
  def flush(): Unit = {
    Await.result(manager.executor.flushdb(), timeout)
  }

  @TearDown
  def shutdown(): Unit = {
    manager.system.terminate()
    Await.result(manager.system.whenTerminated, timeout)
  }
}
```

Result analysis

- Simple charts for JMH benchmarks
<http://nilskp.github.io/jmh-charts/>
- GUI for comparing JMH results
<https://github.com/akarnokd/jmh-compare-gui>



Java vs Scala

divided we fail

Resources

- GC
 - <https://plumbr.eu/java-garbage-collection-handbook>
 - <http://mechanical-sympathy.blogspot.com/2013/07/java-garbage-collection-distilled.html>
- Tools
 - <https://github.com/AdoptOpenJDK/jitwatch>
 - <https://github.com/giltene/jHiccup>
 - <https://github.com/chewiebug/GCViewer>
 - <http://gceasy.io/>
- Microbenchmarking
 - [Nanotrusting the Nanotime](#)
 - [Scala JMH Samples](#)
 - [Java JMH Samples](#)
 - <https://www.youtube.com/watch?v=VaWgOCDBxYw>
- Performance analysis
 - <https://www.youtube.com/watch?v=dqg0R3gYGac>
 - <https://www.youtube.com/watch?v=nZfNehCzGdw>
 - <http://shipilev.net/>
 - <http://mechanical-sympathy.blogspot.com/>

Conclusions

- intuition is almost always wrong (unless you rock)
- never trust anything (unless checked before... and after)
- challenge everything (especially these slides)
- embrace failure (especially your failures)
- grind your teeth, and redo the tests (especially yours)



Thank you!

You can go back to coding now

unless you don't know what to code,

in that case visit: rozkmina.com